

Machine Learning

Exercise 4

Marc Toussaint

Machine Learning & Robotics lab, U Stuttgart
Universitätsstraße 38, 70569 Stuttgart, Germany

May 15, 2017

1 Discriminative Function in Logistic Regression

Logistic Regression (ML 03 Part 1, slides 7 + 8) defines class probabilities as proportional to the exponential of a discriminative function:

$$P(y|x) = \frac{\exp f(x, y)}{\sum_{y'} \exp f(x, y')}$$

Prove that, in the binary classification case, you can assume $f(x, 0) = 0$ without loss of generality.

This results in

$$P(y = 1|x) = \frac{\exp f(x, 1)}{1 + \exp f(x, 1)} = \sigma(f(x, 1)).$$

(Hint: first assume $f(x, y) = \phi(x, y)^\top \beta$, and then define a new discriminative function f' as a function of the old one, such that $f'(x, 0) = 0$ and for which $P(y|x)$ maintains the same expressibility.)

2 Logistic Regression

On the course webpage there is a data set `data2Class.txt` for a binary classification problem. Each line contains a data entry (x, y) with $x \in \mathbb{R}^2$ and $y \in \{0, 1\}$.

a) Compute the optimal parameters β (perhaps also the mean neg-log-likelihood, $-\frac{1}{n} \log L(\beta)$) of logistic regression using linear features. Plot the probability $P(y = 1 | x)$ over a 2D grid of test points. Tips:

- Recall the objective function, and its gradient and Hessian that we derived in the last exercise:

$$L(\beta) = - \sum_{i=1}^n \log P(y_i | x_i) + \lambda \|\beta\|^2 \quad (1)$$

$$= - \sum_{i=1}^n \left[y_i \log p_i + (1 - y_i) \log[1 - p_i] \right] + \lambda \|\beta\|^2 \quad (2)$$

$$\nabla L(\beta) = \frac{\partial L(\beta)}{\partial \beta} = \sum_{i=1}^n (p_i - y_i) \phi(x_i) + 2\lambda I \beta = X^\top (p - y) + 2\lambda I \beta \quad (3)$$

$$\nabla^2 L(\beta) = \frac{\partial^2 L(\beta)}{\partial \beta^2} = \sum_{i=1}^n p_i (1 - p_i) \phi(x_i) \phi(x_i)^\top + 2\lambda I = X^\top W X + 2\lambda I \quad (4)$$

$$\text{where } p(x) := P(y=1 | x) = \sigma(\phi(x)^\top \beta), \quad p_i := p(x_i), \quad W := \text{diag}(p \circ (1 - p)) \quad (5)$$

- Setting the gradient equal to zero can't be done analytically. Instead, optimal parameters can quickly be found by iterating Newton steps: For this, initialize $\beta = 0$ and iterate

$$\beta \leftarrow \beta - (\nabla^2 L(\beta))^{-1} \nabla L(\beta). \quad (6)$$

You usually need to iterate only a few times (~ 10) til convergence.

- As you did for regression, plot the discriminative function $f(x) = \phi(x)^\top \beta$ or the class probability function $p(x) = \sigma(f(x))$ over a grid.

Useful gnuplot commands:

```
splot [-2:3][-2:3][-3:3.5] 'model' matrix \
    us ($1/20-2):($2/20-2):3 with lines notitle
plot [-2:3][-2:3] 'data2Class.txt' \
    us 1:2:3 with points pt 2 lc variable title 'train'
```

b) Compute and plot the same for quadratic features.